

A short summary of the available families is given in the following paragraphs:

`AdaExp()`, `Binomial()` and `AUC()` implement families for binary classification. `AdaExp()` uses the exponential loss, which essentially leads to the AdaBoost algorithm of Freund and Schapire (1996). `Binomial()` implements the negative binomial log-likelihood of a logistic regression model as loss function. Thus, using `Binomial` family closely corresponds to fitting a logistic model. Alternative link functions can be specified.

However, the coefficients resulting from boosting with family `Binomial(link = "logit")` are 1/2 of the coefficients of a logit model obtained via `glm`. Buehlmann and Hothorn (2007) argue that the family `Binomial` is the preferred choice for binary classification. For binary classification problems the response y has to be a factor. Internally y is re-coded to -1 and $+1$ (Buehlmann and Hothorn 2007).

`Binomial(type = "glm")` is an alternative to `Binomial()` leading to coefficients of the same size as coefficients from a classical logit model via `glm`. Additionally, it works not only with a two-level factor but also with a two-column matrix containing the number of successes and number of failures (again, similar to `glm`).

`AUC()` uses $1 - AUC(y, \hat{f})$ as the loss function. The area under the ROC curve (AUC) is defined as $AUC = (n - 1) \sum_{i: y_i = 1} \sum_{j: y_j = -1} I(\hat{f}_i > \hat{f}_j)$. Since this is not differentiable in \hat{f} , we approximate the jump function $I(\hat{f}_i - \hat{f}_j > 0)$ by the distribution function of the triangular distribution on $[-1, 1]$ with mean 0, similar to the logistic distribution approximation used in Ma and Huang (2005).

`Gaussian()` is the default family in `mboost`. It implements `L2Boosting` for continuous response. Note that families `GaussReg()` and `GaussClass()` (for regression and classification) are deprecated now. `Huber()` implements a robust version for boosting with continuous response, where the Huber-loss is used. `Laplace()` implements another strategy for continuous outcomes and uses the `L1-loss` instead of the `L2-loss` as used by `Gaussian()`.

`Poisson()` implements a family for fitting count data with boosting methods. The implemented loss function is the negative Poisson log-likelihood. Note that the natural link function $\log(\mu) = \eta$ is assumed. The default step-size `nu = 0.1` is probably too large for this family (leading to infinite residuals) and smaller values are more appropriate.

`GammaReg()` implements a family for fitting nonnegative response variables. The implemented loss function is the negative Gamma log-likelihood with logarithmic link function (instead of the natural link).

`CoxPH()` implements the negative partial log-likelihood for Cox models. Hence, survival models can be boosted using this family.

`QuantReg()` implements boosting for quantile regression, which is introduced in Fenske et al. (2009). `ExpectReg` works in analogy, only for expectiles, which were introduced to regression by Newey and Powell (1987).

Families with an additional scale parameter can be used for fitting models as well: `PropOdds()` leads to proportional odds models for ordinal outcome variables (Schmid et al., 2011). When using this family, an ordered set of threshold parameters is re-estimated in each boosting iteration. An example is given below which also shows how to obtain the thresholds. `NBinomial()` leads to regression models with a negative binomial conditional distribution of the response. `Weibull()`, `Loglog()`, and `Lognormal()` implement the negative log-likelihood functions of accelerated failure time models with Weibull, log-logistic, and lognormal distributed outcomes, respectively. Hence, parametric survival

models can be boosted using these families. For details see Schmid and Hothorn (2008) and Schmid et al. (2010).

`Gehan()` implements [rank-based estimation](#) of survival data in an [accelerated failure time model](#). The [loss function](#) is defined as the sum of the [pairwise absolute differences of residuals](#). The response needs to be defined as `Surv(y, delta)`, where `y` is the observed survival time (subject to censoring) and `delta` is the non-censoring indicator (see [Surv](#) for details). For details on `Gehan()` see Johnson and Long (2011).

`Cindex()` optimizes [the concordance-index](#) for [survival data](#) (often denoted as Harrell's C or C-index). The concordance index [evaluates the rank-based concordance probability between the model and the outcome](#). [The C-index measures whether large values of the model are associated with short survival times and vice versa](#). The interpretation is similar to the AUC: [A C-index of 1 represents a perfect discrimination while a C-index of 0.5 will be achieved by a completely non-informative marker](#). The `Cindex()` family is based on an estimator by Uno et al. (2011), which incorporates inverse probability of censoring weighting `ipcw`. To make the estimator differentiable, sigmoid functions are applied; the corresponding smoothness can be controlled via `sigma`. For details on `Cindex()` see Mayr and Schmid (2014).

Hurdle models for zero-inflated count data can be fitted by using a combination of the `Binomial()` and `Hurdle()` families. While the `Binomial()` family allows for fitting the zero-generating process of the Hurdle model, `Hurdle()` fits a negative binomial regression model to the non-zero counts. Note that the specification of the Hurdle model allows for using `Binomial()` and `Hurdle()` independently of each other.

Linear or additive multinomial logit models can be fitted using `Multinomial()`; although this family requires some extra effort for model specification (see example). More specifically, the predictor must be in the form of a linear array model (see [%O%](#)). Note that this family does not work with tree-based base-learners at the moment. The class corresponding to the last level of the factor coding of the response is used as reference class.

`RCG()` implements the ratio of correlated gammas (RCG) model proposed by Weinhold et al. (2016).